



Recording your AI use: student guidance

It is expected that students will use AI and LLMs widely on assignments.

This semester, Question 5 requires you to log the actual conversations you have with AI tools, including both your prompts and the AI's replies. The rubric assesses the alignment between your work, your explanation (in a recorded video), and those conversations.

This note covers:

- what to set up before you start;
- use of LLMs;
- what to log; and
- how to save and submit your logs.

Before you start

Set this up in your first week, before you begin work.

Every way students lose logs is retrospective. History gets purged, a temporary chat expires, a folder gets cleared. Nothing you do at the end of the assignment will recover a conversation you had in week two, so the habit of saving each conversation when you finish with it is the single thing that makes this requirement easy rather than painful.

Create your `ai_logs` folder now, and take the one-off step for whichever tools you expect to use:

- **Chat assistants** (ChatGPT, Claude, Gemini, Copilot Chat). Save each conversation to a file when you finish with it. Do not rely on the tool's own history, and do not use temporary or incognito chats for assignment work, as those cannot be recovered.
- **Claude Code**. Run `/export` at the end of each session. Sessions are also stored locally, but are deleted after 30 days by default, so raise `cleanupPeriodDays` in your settings now as a backstop rather than something to rely on.
- **Codex CLI**. Codex has no export command. Sessions are saved automatically as `.jsonl` files under `~/codex/sessions/`. Copy the session file into your `ai_logs` folder at the end of each session, named as usual. The compile script converts these for you, so nothing further is needed.
- **Work accounts**. Check that copying or exporting out of the tool is permitted by your employer, and is not blocked, before you rely on it.



Then capture one conversation from beginning to end and open the saved file to check the whole exchange is there. If something in your setup does not work, it is much better to discover it in week one than the night before submission.

Use of LLMs

Corporate vs personal accounts

We recommend using a personal account rather than one provided by your employer. Work accounts may block exporting chats, be monitored or otherwise managed by your employer, and can mix confidential work data into your logs — which you must not submit.

Free vs frontier models

The assignment is scoped at approximately 50 hours. It can be completed with free tiers, or entirely without AI, but students taking those routes should budget more time. With command-line coding agents (Claude Code, Codex CLI), it can be completed in under 50 hours.

Free tiers have usage caps that a project of this size will exhaust quickly, and offer a limited range of models. Frontier models and coding agents generally require a paid plan or API credits. Many students will choose a personal paid subscription for the assignment period; a single entry-level subscription is typically sufficient, and no specific vendor or plan is required or endorsed.

Other considerations

In practice, actuaries increasingly work with these tools, and the assignment is designed to reflect that working environment.

If a work account is genuinely your only option: check your employer permits it, keep employer and client data out of the conversations, and record any capture limitations honestly (see *If capture fails*, below).



What to log

Overview

Keep every conversation you have with an AI: a chat assistant, an editor chat, a command-line agent, or a model you call from your own code. Keep the whole exchange, both your messages and the replies, as the tool gives it to you. Do not summarise it and do not write it up afterwards. There is nothing to keep where there was no conversation, such as inline autocomplete or work you did yourself.

Submit your logs as exported. Do not tidy, edit, or delete anything yourself — an unedited log, including the dead ends and the things that did not work, is worth more than a curated one. The compile script may parse command-line agent sessions to leave out tool output and internal reasoning, so do not be surprised if those are absent from the compiled file.

The amount of AI you use never directly affects your marks. What Question 5 marks is the alignment between your work, your video explanation, and your logged conversations. Markers sample your conversations as supporting evidence; nobody reads the whole log, and nothing is marked for volume, tidiness or completeness.

An occasional honest gap, recorded in your log, will not reduce marks, but major gaps will limit the marks available to you. Falsifying a record, or denying AI use that occurred, is an academic integrity matter, separate from marking.

Record as you go. Save each conversation as its own dated file on the day you finish with it.

How to keep a conversation

The general approach is to select the whole conversation, both your messages and the AI's replies, and copy-paste it into a text or markdown file. The tool-specific shortcuts below are optional, but may be easier in some cases.



How you're using AI	How to log it
A chat assistant (ChatGPT, Claude, Gemini, Copilot Chat)	Copy-paste the full exchange into a file. If your tool has a proper export to file, that is fine too. Do not submit share links , as they can break, be deleted, or miss later messages; only a saved copy of the text counts as your record.
A coding assistant you chat with in your editor (Copilot Chat, Cursor chat)	Same: copy the whole thread into a file. Many editor chats have no export, so copy-paste is the reliable route.
A command-line coding agent (Claude Code, Codex CLI)	In Claude Code, run /export at the end of each session and save the markdown file into ai_logs. Sessions are also stored under ~/.claude/projects/, but as .jsonl files that are deleted after 30 days by default, so export as you go. Codex has no export command: copy the session's .jsonl file from ~/.codex/sessions/ into ai_logs, and the compile script converts it for you. Before submitting any agent transcript, skim it for API keys, credentials, or confidential work data and remove them.
Calling an AI model from code (an API call in your notebook)	Use the provided log_ai() helper in the template to record each prompt and the model's full response to a file as you run, or keep your own save of both.

If a tool is not listed, use the closest row or ask in the discussion forum.

For rich content (canvases, generated images, uploaded files, research reports), capture the text of the exchange. Only text can be compiled into what you submit, so if the tool produced something that does not copy-paste, add a line to the log file at the point it was produced saying what it was.

If capture fails

Sometimes you cannot keep a conversation: a work tool blocks export, history was deleted, a temporary chat expired, or you simply forgot. The same applies if you lose your logs entirely, for example if your files are wiped.



Before concluding that anything is lost, check the tool's own history, as most chat assistants keep your conversations on their servers even when your local files are gone.

None of this is a breach, provided you record it honestly in a manual log entry: a short file of your own, saved in `ai_logs` and named like any other (the Getting Started notebook has example code for writing one). The marking and integrity consequences of gaps are set out in the Overview above.

Saving and submitting your logs

Create an `ai_logs` folder (on your local machine or a cloud drive). Save each conversation as its own file in that folder, named with the date first:

```
ai_logs/  
2026-08-03_Q1_chatgpt_tfidf.md  
2026-08-05_Q1_copilot_plotting.txt  
2026-08-09_multiQ_claude_data_cleaning.md  
2026-08-14_Q3_claudecode_session.md  
2026-08-20_general_claude_glm_reading.md
```

Name each file using these rules, so that the provided script can read it:

- date first (YYYY-MM-DD), the day of the conversation, so files sort chronologically. Optionally add a time (2026-08-05_1430_Q1_...) if you have several conversations the same day and want them ordered; do not guess a time you do not know, as the date alone is fine;
- then the question tag: Q1–Q4, or multiQ for a conversation spanning questions, or general for one from before you had picked a question, both of which are normal; and
- then a couple of words on what it was about.

At the end, run the provided script from your submission folder:

```
python compile_ai_logs.py
```

It generates **`ai_logs/AI_LOGS.md`**: a dated contents table followed by every conversation in chronological order. It never changes or deletes your saved conversations, and you can re-run it at any time.

The script also flags anything it could not read, could not convert, or that does not follow the naming pattern. Read those warnings before you submit: a conversation the script could not read is a conversation your marker will not see.



DSA S2 2026 Assignment

Recording your AI use

Submit AI_LOGS.md. It is the only file submitted, so keep your ai_logs folder as your own record but do not rely on anything in it reaching your marker. Do not edit AI_LOGS.md after generating it — if something in it is wrong, fix the underlying file and re-run the script.